

# Mark Schott

Senior Cloud Data Engineer

Old Lyme, CT · 313-908-0388 · hire.mark.schott@gmail.com  
linkedin.com/in/meschott · github.com/meschott

Senior Cloud Data Engineer with 8+ years building production data platforms on AWS, drawing on a physics + applied-statistics foundation for domain modeling. Recently led greenfield AWS + DBT/Snowflake migration at EDF Power Solutions; prior work powered KPIs for 80+ solar/storage sites across 6 countries. Daily LLM-assisted development across cloud and local models.

## SKILLS

### Languages

Python (uv, pre-commit), SQL, Bash

### AWS / Data

Lambda, Glue, EMR, Athena, S3, DynamoDB, RDS Aurora, PySpark, Pandas, Hudi

### DevOps / IaC

AWS CDK, Terraform, CodePipeline, GitHub Actions, Jenkins, Git, Docker (Buildx x86/ARM), DBT

### Analytics

Time-series, outlier detection, regression; matplotlib, Altair + Polars, Scikit-learn, NetworkX

### Other

Linux, Snowflake, Jira

## AI TOOLING

OpenCode, GitHub Copilot · Anthropic & OpenAI APIs · local Qwen / Gemma · plan / build agent workflows

## EDUCATION

### BS Physics, *cum laude*

Wayne State University · 2013

### Applied Physics graduate coursework

Rice University · 2013–14

### Data Science Certificate

NYC Data Science Academy · 2017

## LANGUAGES

English · French (Advanced)

## EXPERIENCE

### EDF Power Solutions Senior Cloud Data Engineer

Dec 2021 – Present

Remote

- **Led 2-engineer team through 6-month corporate platform migration:** stood up new AWS account (CDK), migrated 80+ DBT models to a new Snowflake instance (Terraform Snowflake provider), designed least-privilege IAM, coordinated across 3 teams; delivered on schedule.
- **Refactored inherited CDK IaC repo** into a clean, deployable baseline; shipped to production AWS via GitHub Actions.
- **Automated Informatica Secure Agent install on Linux** powering 16 auto-ingest Snowpipes (AWS → Snowflake).
- **Built a JQL visualizer** (Python · Jira Cloud API · NetworkX) to map cross-team dependencies and surface blockers during the migration.
- **Refactored 10+ AWS Glue jobs** from single- to multi-partition processing — cut compute cost by **over 100×**; an additional **\$10K+/year** recovered via root-cause analysis on pipeline waste and data quality.
- Built KPI suite (Potential Production, O&M, snow loss) on AWS Lambda + Glue + Athena for **80+ solar/storage sites across 6 countries**.

### Fiserv Data Scientist

Jan – Nov 2021

New Providence, NJ

- Dockerized Python app for Kubernetes; became primary developer and code owner.
- Built and packaged a Cython interface to an external C library for column-level encryption.
- Drove Jira / Agile adoption and authored team documentation; demonstrated a local-dev workflow that eliminated need for cloud VMs.
- Integrated Azure ADLS + Databricks for cross-cloud handoff.

### ADP Senior Data Engineer (promoted from Data Engineer)

Sep 2017 – Nov 2020

Roseland & Parsippany, NJ

- SME and developer for a business-critical PySpark pipeline driving **\$100M+ in CARES loans** and a multi-million-dollar data-monetization business.
- Refactored outsourced product to in-house PySpark on EMR + Glue + S3; **cut client data-refresh time by 2+ weeks**.
- Built time-series outlier detection (Python / Scala) catching pipeline anomalies in-flight, saving hours of research and ops per incident.
- Defined EU-compliant data-anonymization standards enabling global benchmarking product rollout.

Earlier: Applications Developer, Reynolds & Reynolds (2015–16) · Experimental Research Intern, Wayne State (2011–13)